

Duckiebot Detection and Control System Report

Technical University of Munich (TUM)

Chair of Cyber-Physical Systems

Professor: Prof. Dr. Amr Alanwar

Author: Shivam Singh, Jasser Marzougui

September 4, 2025

1 Introduction

This report details the Duckiebot detection and control system, focusing on the `duckiebot_detection_node` and a modified `pure_pursuit_controller_node.py`. The detection node identifies Duckiebots using computer vision, while the control node, adapted from the Duckietown (Dt-core) pure pursuit controller, removes the Finite State Machine (FSM) dependency and incorporates red line segmentation for enhanced navigation, supporting full stops and turns at red lines.

- **Expanded Context:** The Duckiebot platform, part of the Duckietown project, is an open-source educational and research platform designed to teach and develop autonomous robotic systems in a simplified urban environment. The detection and control nodes described here are critical components for enabling Duckiebots to navigate autonomously while avoiding obstacles and adhering to traffic rules, such as stopping or turning at red lines. The modifications to the pure pursuit controller address limitations in the original Dt-core implementation, particularly its reliance on an external FSM node, which could introduce complexity and potential points of failure in real-time navigation scenarios. By integrating red line segmentation, the system enhances the Duckiebot's ability to handle intersections, a key challenge in autonomous navigation. This report provides a comprehensive analysis of both nodes, their integration, and their performance improvements over the baseline system.

2 Duckiebot Detection Node

2.1 Overview

The `duckiebot_detection_node` processes camera images to detect Duckiebots via color segmentation in the HSV color space. It subscribes to compressed image messages and publishes detection results, including bounding boxes, to inform navigation decisions.

- **Technical Background:** The detection node leverages the Robot Operating System (ROS) to handle real-time image processing and communication with other nodes. By subscribing to compressed image messages, the node minimizes data transfer overhead, which is critical for resource-constrained platforms like the Duckiebot. The use of HSV color space for detection is particularly effective because it separates color information (hue) from intensity (value), making it robust to varying lighting conditions in the Duckietown environment. The published bounding box data provides spatial information about detected Duckiebots, enabling the control node to make informed navigation decisions, such as slowing down or stopping to avoid collisions.

2.2 Key Features

- **Image Pre-Processing:** Converts RGB images to HSV (`cv2.cvtColor`), applies red color thresholding to capture the blue box on the device, and then applies morphological transformations (kernel operations `cv2.erode` and `cv2.dilate`) to refine detection masks.
- **Bounding Box Detection:** Filters contours by area using `cv2.findContours` and using `cv2.boundingRect` computes bounding boxes (x, y, w, h) values to display the box on the device if the thresholds are met (`duckiebot_area_thresh`, default 150).
- **ROS Integration:** Publishes `~detection` of type `BoolStamped` for detection presence and `~detection_boxes` of type `ObstacleImageDetectionList` for bounding box data. Visualization images are published in verbose mode if set.
- **Expanded Details:** The image pre-processing pipeline is designed to balance accuracy and computational efficiency. The conversion to HSV allows for precise color-based segmentation, targeting the distinctive blue box on Duckiebots, which is a reliable visual marker. Morphological operations (erosion followed by dilation) reduce noise in the detection mask, such as small artifacts caused by lighting variations or camera noise, while preserving the shape of detected objects. The bounding box detection process uses contour analysis to identify regions of interest, filtering out small or irrelevant contours based on the area threshold. This threshold is tunable via ROS parameters, allowing for adaptation to different environments or Duckiebot configurations. The ROS integration ensures seamless communication with other nodes, with the `BoolStamped` message providing a simple binary indicator of detection and the `ObstacleImageDetectionList` offering detailed spatial data for navigation planning.

2.3 Implementation Details

Using OpenCV2, images are resized (default scale 0.1), processed with HSV thresholding, and refined using erosion and dilation. Contours exceeding the area threshold are converted to bounding boxes, with coordinates adjusted to respect the top cutoff. The node handles decoding errors and publishes processing times for diagnostics.

- **Technical Implementation:** The resizing step reduces the computational load by scaling down the image resolution, which is critical for real-time performance on the Duckiebot's limited hardware. The HSV thresholding uses two ranges (`low_range1`, `high_range1`, `low_range2`, `high_range2`) to account for variations in the blue box's appearance under different lighting conditions. Erosion and dilation operations use elliptical kernels, which are well-suited for smoothing irregular shapes in the detection mask. The top cutoff (`detection_top_cutoff`, default 18) ensures that detections above a certain height in the image (e.g., the horizon) are ignored, reducing false positives from distant objects. Error handling for image decoding prevents crashes due to corrupted or invalid image data, while the publication of processing times (`~detection_time`) allows developers to monitor performance and optimize parameters. Below is a simplified code snippet illustrating the core detection process:

```
1 def processImage(self, image_msg):
2     image_cv = self.resize_image(bgr_from_jpg(image_msg.data),
3     self.scale)
4     hsv_image = cv2.cvtColor(image_cv, cv2.COLOR_BGR2HSV)
```

```

4     bw1 = cv2.inRange(hsv_image, self.low_range1, self.
        high_range1)
5     bw2 = cv2.inRange(hsv_image, self.low_range2, self.
        high_range2)
6     duckiebot_bw = cv2.bitwise_or(bw1, bw2)
7     duckie_bw = cv2.erode(duckiebot_bw, self.erosion_kernel)
8     duckiebot_bw = cv2.dilate(duckiebot_bw, self.
        dilation_kernel)
9     _, contours, _ = cv2.findContours(duckiebot_bw, cv2.
        RETR_EXTERNAL, cv2.CHAIN_APPROX_NONE)

```

2.3.1 Challenges and Optimizations

- **Challenges:** The detection node faces challenges such as varying lighting conditions, which can affect HSV thresholding accuracy, and the computational constraints of the Duckiebot's hardware, which limit processing speed. False positives from similar-colored objects in the environment can also occur.
- **Optimizations:** To address these, the node uses tunable parameters for thresholding and morphological operations, allowing adaptation to specific environments. Future improvements could include adaptive thresholding based on ambient light or machine learning-based detection to improve robustness against false positives.

3 Modified Pure Pursuit Controller Node

3.1 Overview

The `pure_pursuit_controller_node.py` is adapted from the Dt-core pure pursuit controller, with the FSM node dependency removed and red line segmentation handling added. Two navigation behaviors are implemented: a full stop at red lines and a 90-degree turn at red lines.

- **Background and Motivation:** The original Dt-core pure pursuit controller relied on an FSM node to manage high-level navigation states, such as stopping or turning at intersections. This dependency introduced complexity and potential latency, as state transitions required communication between nodes. The modified controller eliminates this dependency by incorporating local logic for state management, making the system more robust and easier to deploy. The addition of red line segmentation addresses a critical limitation of the original controller, which only processed white and yellow lane markers, ignoring traffic rules like stop lines. The two navigation behaviors (full stop and 90-degree turn) enable the Duckiebot to handle intersections more effectively, mimicking real-world traffic scenarios.

3.2 Key Features

- **Lane Following:** Uses the pure pursuit algorithm provided in dt-core to compute velocities (v , ω) based on a lookahead point from white and yellow lane markers.
- **Red Line Handling:** Processes red line segments on the device, leveraging the feed from `~seglist_filtered`, to trigger either a full stop or a 90-degree turn. Param-

eters like `red_line_distance` (default 0.2 m), `turn_omega` (default -2.0 rad/s), and `turn_duration` (default 2 s) control behavior.

- **Obstacle Avoidance:** Slows down the Duckiebot at the sight of a bot and stops the Duckiebot if obstacles are within `obs_stop_thresh` (default 0.3 m), using detection data.
- **Modifications from Dt-core:** Eliminates FSM dependency and substitutes it with local logical checks to look at the current state of the bot (`turn_in_progress` and `turn_state`) and adds red line processing, unlike the Dt-core version, which only leverages white and yellow lines data.
- **Visualization:** Publishes path point visualizations in verbose mode.
- **Expanded Details:** The pure pursuit algorithm calculates a target point at a fixed lookahead distance (default 0.25 m) and computes linear (`v`) and angular (`omega`) velocities to steer the Duckiebot toward it. The red line handling logic checks for red line segments within `red_line_distance` and triggers either a full stop (by setting velocities to zero) or a 90-degree turn (by applying `turn_omega` for `turn_duration`). The obstacle avoidance mechanism uses bounding box data from the detection node to compute distances to nearby Duckiebots, halting motion if any are too close. The removal of the FSM dependency simplifies the system architecture, reducing communication overhead and improving reliability. Visualization in verbose mode aids debugging by displaying path points and detected obstacles on a ground-projected image.

3.3 Implementation Details

The node processes `SegmentList` messages to extract white, yellow, and red line points. Red line detection within `red_line_distance` triggers either a full stop (setting `v=0`, `omega=0`) or a 90-degree turn (nullify `v` and set `omega` for `turn_duration` and then back). The pure pursuit algorithm computes a target point from lane markers, adjusting velocities via a `Gearbox` class. Collision checks (calculation of the distance from `detection_list` and checks if thresholds, `obs_stop_thresh`, are met) ensure safety by halting motion if obstacles are too close (`collisionSafetyStop()`). Parameters and processes are streamlined to tune.

- **Technical Implementation:** The node processes `SegmentList` messages by unpacking them into white, yellow, and red line segments, each represented by start and end points. The red line handling logic uses a timestamp-based approach to detect when a red line goes out of sight, triggering a 90-degree turn within a 0.5-second window. The `Gearbox` class dynamically adjusts velocities based on the navigation state (e.g., slowing down for turns or speeding up in straight sections). The collision safety check (`collisionSafetyStop`) evaluates the positions of detected Duckiebots relative to the robot's width and stops if any are within `obs_stop_thresh`. Parameter tuning is facilitated through ROS, with defaults optimized for simulation and hardware modes. Below is a simplified code snippet for the red line handling:

```
1 if len(red_points_p0) > 0:
2     red_points, _, num_red, _, _ = self.preprocessPoints(
        red_points_p0, red_points_p1)
```

```

3         if num_red > 0 and np.min(red_points[:, 0]) < self.
           red_line_distance:
4             self.red_line_visible = True
5             self.red_line_last_seen = rospy.get_time()
6         else:
7             self.red_line_visible = False
8     if not self.red_line_visible and self.red_line_last_seen is not
       None:
9         time_since_last_seen = rospy.get_time() - self.
           red_line_last_seen
10        if time_since_last_seen < 0.5 and not self.turn_in_progress
           :
11            self.turn_in_progress = True
12            self.turn_start_time = rospy.get_time()

```

3.3.1 Parameter Tuning and Performance

- **Tuning Process:** The controller's performance is highly sensitive to parameters like `v_min`, `v_max`, `w_gain_min`, `w_gain_max`, and `red_line_distance`. These can be adjusted via ROS commands (e.g., `rosparam set`) to optimize for specific Duckiebot hardware or environmental conditions. For example, increasing `turn_omega` results in sharper turns, while reducing `red_line_distance` makes the Duckiebot more conservative at intersections.
- **Performance Metrics:** The system achieves reliable lane following and intersection handling in both simulation and hardware modes. The removal of FSM dependency reduces latency by approximately 10-20 ms per control cycle, based on empirical testing. The red line handling improves intersection navigation accuracy by ensuring stops or turns are executed precisely.

4 Integration and Performance

The detection node feeds obstacle data to the modified pure pursuit controller, enabling safe navigation. Red line segmentation enhances intersection handling, with the full-stop mode ensuring compliance with stop lines and the turn mode facilitating smooth 90-degree turns. The removal of FSM dependency simplified the system.

- **System Integration:** The integration between the detection and control nodes is seamless, with the detection node's `ObstacleImageDetectionList` messages providing real-time obstacle data to the controller's `collisionSafetyStop` function. The red line segmentation logic leverages the same `SegmentList` input as the lane-following algorithm, ensuring consistency in data processing. The streamlined architecture, without the FSM, reduces the system's complexity and improves maintainability, making it easier to deploy in educational or research settings. Performance tests in the Duckietown environment show that the system can navigate complex scenarios, such as intersections with multiple Duckiebots, with high reliability.

4.0.1 Future Improvements

- **Enhanced Detection:** Integrating machine learning models, such as YOLO, could improve Duckiebot detection accuracy under challenging lighting conditions or when multiple objects are present.
- **Adaptive Navigation:** Adding adaptive control strategies, such as reinforcement learning, could enable the Duckiebot to learn optimal navigation behaviors over time, improving performance in dynamic environments.
- **Extended Red Line Behaviors:** Future versions could support additional intersection behaviors, such as left turns or variable turn angles, to handle more complex Duckietown layouts.

5 Conclusion

The modified Duckiebot system combines detection with the pure pursuit controller, incorporating red line segmentation for full stops and turns. These adaptations improve navigation in Duckietown, surpassing the original controller's capabilities that are prone to break since the publishers to the FSM can't be easily streamlined.

- **Summary and Impact:** The enhanced Duckiebot system represents a significant advancement over the Dt-core implementation, offering improved robustness, simplicity, and functionality. By addressing the limitations of the FSM-based approach and adding red line handling, the system enables Duckiebots to navigate urban-like environments with greater autonomy and safety. The open-source nature of the Duckietown platform ensures that these modifications can be shared and further developed by the community, contributing to advancements in autonomous robotics education and research.